

ENEL464 Embedded Software and Advanced Computing 2024

Group Assignment 2

1 Introduction

The purpose of this assignment is to implement a numerical algorithm on a computer to make the most efficient use of its caches and cores. You will find that you will get a large variation in program performance depending on how you implement your code and use compiler optimisations. A knowledge of computer architecture should help with this!

The algorithm is Jacobi relaxation. This is an iterative algorithm used to approximate differential equations, for example, Poisson's and Laplace's equations. Poisson's equation can be used to find the electric potential given a specified charge distribution or temperature given a specified heat source. There are more efficient ways to solve this problem using Green's functions and Fourier transforms but that is not the purpose of this assignment.

2 Jacobi relaxation

The discrete form of Poisson's equation is

$$\nabla^2 V_{i,j,k,n} = f_{i,j,k,n}, \quad (1)$$

where $f_{i,j,k,n}$ is the source for a voxel with coordinates i, j, k at timestep n and $V_{i,j,k,n}$ is the potential field to be determined for the same voxel. If f is the source charge, V is the electric potential.

Poisson's equation is a partial differential equation and thus requires boundary conditions to achieve a solution. For the assignment, the potential on the bottom boundary ($k = 0$) is -1 V, the potential on the top boundary ($k = N - 1$) is 1 V, and the four sides are insulated. This is a mixed boundary condition problem since the bottom and top sides have Dirichlet boundary conditions and the sides have Neumann boundary conditions.

Poisson's equation can be solved iteratively, at each time-step n , using Jacobi relaxation, where¹

$$V_{i,j,k,n+1} = \frac{1}{6} \left(V_{i+1,j,k,n} + V_{i-1,j,k,n} + V_{i,j+1,k,n} + V_{i,j-1,k,n} + V_{i,j,k+1,n} + V_{i,j,k-1,n} - \Delta^2 f_{i,j,k} \right). \quad (2)$$

¹You might recognise (2) as a 3-D discrete convolution at each time-step.

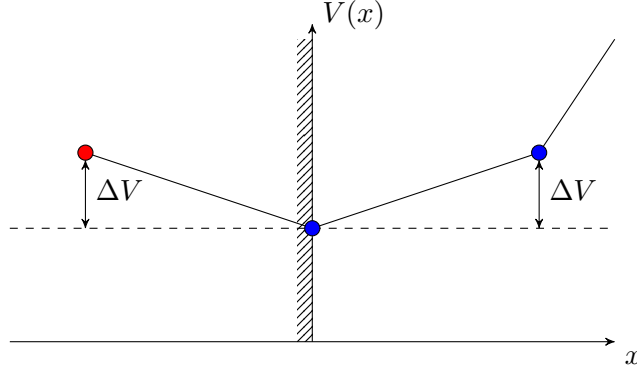


Figure 1: Illustration of a Neumann boundary condition where the external ‘ghost’ point (red) is defined such that the gradient at the boundary is zero with respect to the internal points (blue).

Here $\Delta = \Delta x = \Delta y = \Delta z$ is the spacing between voxels in metres, $0 \leq i < N$, $0 \leq j < N$, and $0 \leq k < N$.

Voxels on the side boundaries ($i = 0$, $i = N - 1$, $j = 0$, $j = N - 1$), are defined such that the derivative over that boundary is zero ($\partial V / \partial n = 0$), see Fig. 1. For example, consider the derivative in the x-direction, $\partial V / \partial x$. This can be approximated by the central difference,

$$\frac{V_{i+1,j,k} - V_{i-1,j,k}}{2\Delta}. \quad (3)$$

On the boundary where $i = 0$, then

$$\frac{V_{1,j,k} - V_{-1,j,k}}{2\Delta}, \quad (4)$$

and for this to be zero implies $V_{-1,j,k} = V_{1,j,k}$. Since these voltages are the same, no current flows normal to the boundary as expected for an insulating boundary. Similarly, on the other x-boundary, $V_{N,j,k} = V_{N-2,j,k}$.

3 Implementation

Your program to solve (2) must either use C or C++. Your goal is to find a fast implementation that will run on your (or an ECE lab) computer, making best use of the caches and multiple cores.

A good starting point is provided with `poisson.c`. You should modify this to solve the assignment. Instructions on building it are in the source code. Some suggested reading is located in the `docs/` folder regarding compiler optimisations, multithreading, and profiling.

Code templates can be found at <https://eng-git.canterbury.ac.nz/mp/ence464-assignment-2024>. We recommend writing your solutions in a Linux-like environment. This means any Linux distribution, MacOS with Homebrew, or Windows with WSL. Your mileage may vary with the profiling tools in anything other than Linux.

4 Testing

Test your algorithm with a single point charge in the centre of the volume, i.e.,

$$f_{i,j,k} = \begin{cases} 1 & i = N/2, j = N/2, k = N/2, \\ 0 & \text{otherwise} \end{cases}, \quad (5)$$

where the volume is comprised of $N \times N \times N$ voxels. Note, your algorithm must work with **arbitrary source distributions**. You cannot assume symmetry.

A testing script (`test.sh`) has been provided along with some sample outputs at several cube sizes. This will automatically compare your output against a reference implementation.

Once a correct solution has been implemented, rename `gitlab-ci.yml` to `.gitlab-ci.yml` to enable continuous integration. The pipeline will run your code against the reference solutions to confirm the correctness of your algorithm. Please do not attempt to use this for performance testing.

5 Support

Only questions submitted via the ENCE464 assignment forum will be answered. Emails will be quietly ignored.

6 FAQ

- *How do I find out the CPU version?* Run `lscpu`. You can also run `lshw` but this needs root privileges to get the cache sizes. Note, Linux considers each thread of a multithreaded core to be a CPU.
- *Why is my program killed?* This is due to the Linux out of memory killer; you have tried to allocate too much memory.

7 Reports

The reports are group reports and are to be submitted as PDF documents through the ENCE464 Learn page. They will be submitted to TurnItIn for plagiarism checking. Use the supplied template and follow the instructions within.

Guidelines for writing a report are available at <https://eng-git.canterbury.ac.nz/mpd/report-guidelines/blob/master/report-guidelines.pdf>. However, no titlepage, abstract, introduction, or conclusion are required.

8 Assessment

Your report will be marked in terms of:

- Written style. The writing should be concise technical writing.
- Presentation. The key here is consistency and clear diagrams and graphs.
- Architecture overview. This should describe your computer's architecture (such as the cache organisation, memory size, CPU model, etc.).
- Multithreading analysis. This should discuss how your program takes advantage of multiple cores.
- Cache analysis. This should discuss how your program takes advantage of the cache.
- Profiling analysis. This should show which parts of your program take the most time.
- Optimisation analysis. This should discuss the affects of some of the compiler optimisations, such as loop unrolling, on your program.
- Individual topic. This should address an aspect of computer architecture related to the assignment. For example, using SIMD instructions, using the GPU, comparing two different computers, analysing branch prediction, optimisation, caching etc. in more detail. The topics must be unique per group.

Five bonus marks will be awarded to any group who can beat our (not very good) program when running on a randomly selected ESL lab computer and give the correct results.

Your report should present the statistics for the time your program takes to run for the following problem sizes: $N = 101, 201, 301, 401, 501, 601, 701, 801, 901$. For $N = 801$ and $N = 901$, do not worry about performing many trials. If you have a really old computer, you can skip $N = 901$.

Your analyses should be backed up with experimental evidence, such as the output from profiling tools.

9 Code

We require you to use git for version control. We will look at your code if we suspect plagiarism.

Your fastest implementation must be an executable named `poisson` that is built by running `make` at the command line. For testing purposes, it must accept two command line flags:

- `-i`: the number of iterations
- `-n`: the edge length of the cube.

It must print out the resulting values from the middle slice of the cube in space-separated format to five decimal places. (You can just use the command line argument parsing and output formatting code from the provided `poisson.c` example.)